

Week 1

# GNU/Linux

- Open-source operating system
  - **Kernel:** core of operating system
    - Allocates time and memory to programs
    - Handles file system and communication between software and hardware
  - **Shell:** interface between user and kernel
    - Interprets commands user types in
    - Takes necessary action to cause commands to be carried out
  - **Programs**
- Why is the system broken up into these different parts?

# Files and Processes

- Everything is either a **process** or a **file**:
  - **Process**: an executing program identified by PID
  - **File**: collection of data
    - A document
    - Text of program written in high-level language
    - Executable
    - Directory
    - Devices

# The Basics: Shell

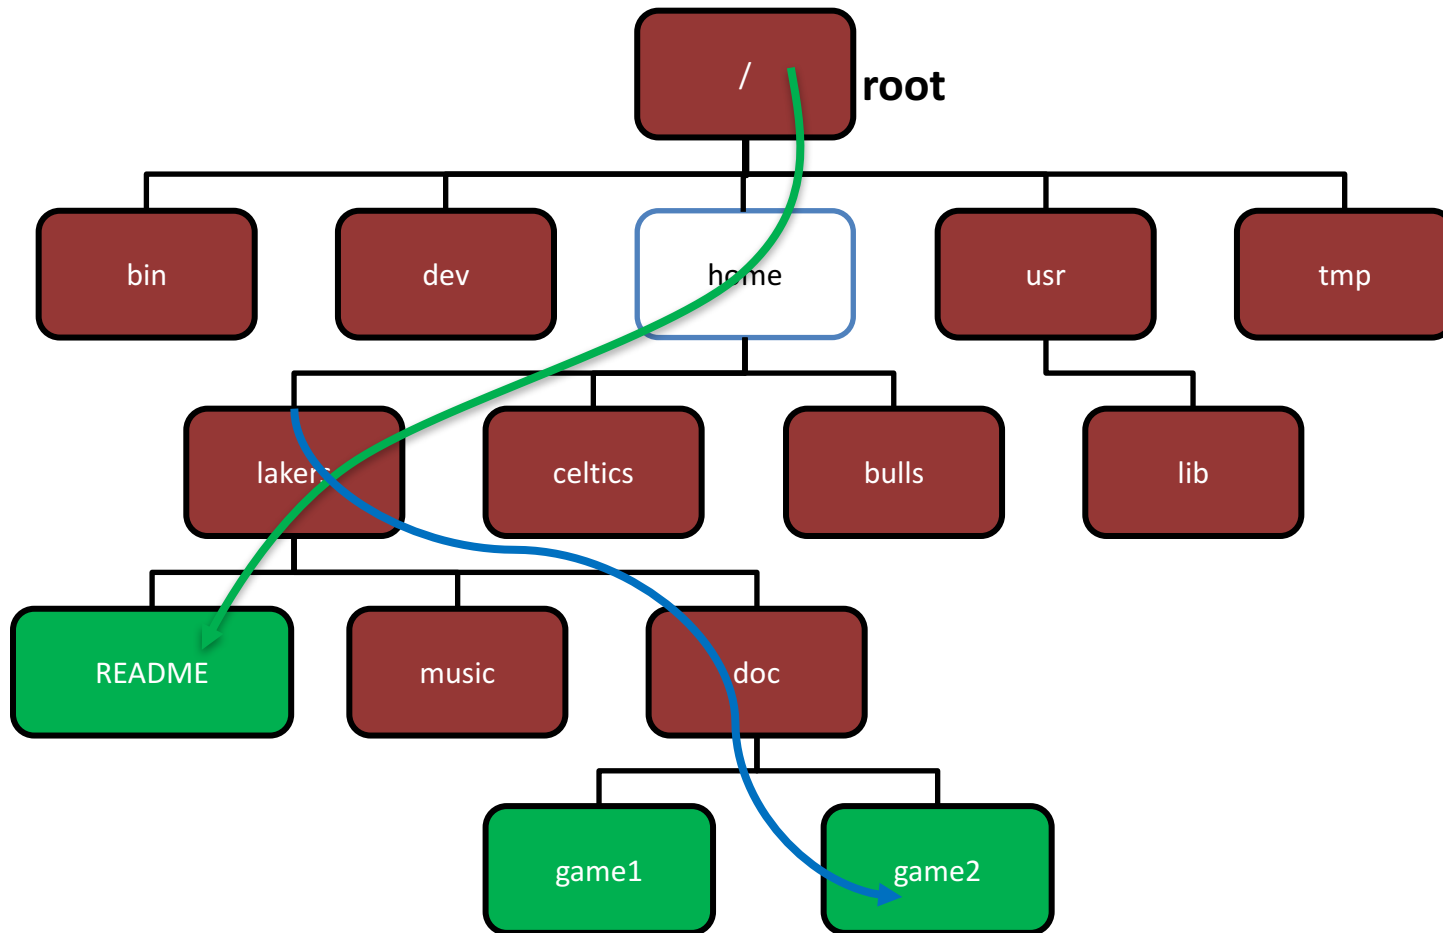
CLI utilities from week 1 you should be familiar with:

- |         |           |        |
|---------|-----------|--------|
| – pwd   | – ls      | – ps   |
| – cd    | – ln      | – kill |
| – mv    | – touch   | – diff |
| – cp    | – find    | – wget |
| – rm    | – whatis  |        |
| – mkdir | – whereis |        |
| – rmdir | – man     |        |

# The Basics: Shell

- How do I find where files are on the system?
- How do I find out what options are available for a particular utility?
- When is a file a file and when is it a process?
- What types of links are there?

# Absolute Path vs. Relative Path



Current directory: home    What are the differences between absolute and relative paths?

# Linux File Permissions

- `chmod`
  - read (r), write (w), executable (x)
  - User, group, others
- Why do we have permissions at all?

Reference	Class	Description
u	user	the owner of the file
g	group	users who are members of the file's group
o	others	users who are not the owner of the file or members of the group
a	all	all three of the above, is the same as <i>ugo</i>

# The Basics: chmod (symbolic)

Operator	Description
+	adds the specified modes to the specified classes
-	removes the specified modes from the specified classes
=	the modes specified are to be made the exact modes for the specified classes

Mode	Name	Description
r	read	read a file or list a directory's contents
w	write	write to a file or directory
x	execute	execute a file or recurse a directory tree



# The Basics: chmod (numeric)

#	Permission
7	full
6	read and write
5	read and execute
4	read only
3	write and execute
2	write only
1	execute only
0	none

- Usage

– `chmod ["references"]["operator"]["modes"] "file1" ...`

Example: **chmod** ug+rw mydir, **chmod** a-w myfile,

Example: **chmod** ug=rx mydir, **chmod** 664 myfile

Week 2

# Locale

## A locale

- Set of parameters that define a user's cultural preferences
  - .Language
  - .Country
  - .Other area-specific things
- What else does the locale affect?

`locale` command

prints information about the current locale environment to standard output

# Environment Variables

- Variables that can be accessed from any child process
- Why do we have these at all? What functions do they serve?

Common ones:

- **HOME**: path to user's home directory
- **PATH**: list of directories to search in for command to execute
- Change value:  
    `export VARIABLE=...`

# Locale Settings Can Affect Program Behavior!!

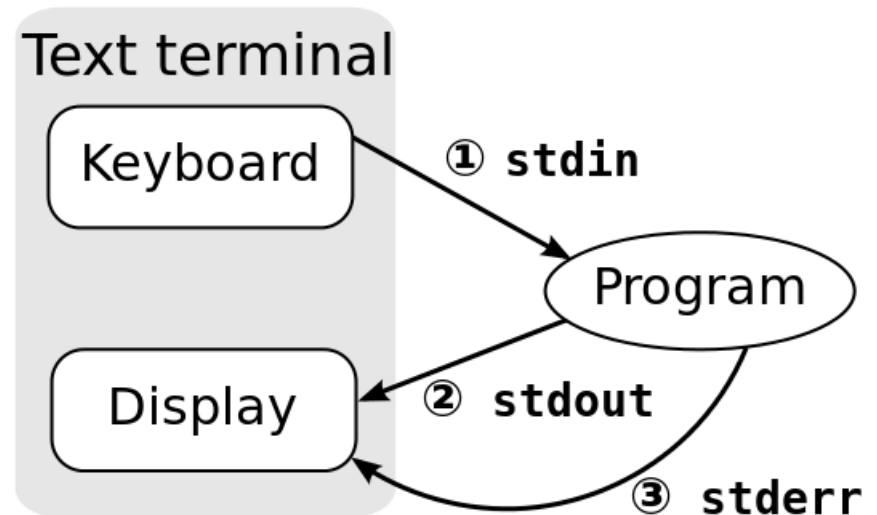
Default sort order for the `sort` command depends:

- `LC_COLLATE='C'`: sorting is in ASCII order
- `LC_COLLATE='en_US'`: sorting is case insensitive except when the two strings are otherwise equal and one has an uppercase letter earlier than the other.

Other locales have other sort orders!

# Standard Streams

- Every program has these 3 streams to interact with the world
  - `stdin` (0): contains data going into a program
  - `stdout` (1): where a program writes its output data
  - `stderr` (2): where a program writes its error msgs



# Redirection and Pipelines

- *program* < *file* redirects *file* to *programs's* *stdin*:  
`cat <file`
- *program* > *file* redirects *program's* *stdout* to *file2*:  
`cat <file >file2`
- *program* 2> *file* redirects *program's* *stderr* to *file2*:  
`cat <file 2>file2`
- *program* >> *file* **appends** *program's* *stdout* to *file*
- *program1* | *program2* assigns *stdout* of *program1* as the *stdin* of *program2*; text '*flows*' through the pipeline  
`cat <file | sort >file2`

Why would we want to redirect I/O? What are some examples of use cases for I/O redirection? How do we implement this in C?

# Regular Expressions

- Notation that lets you search for text with a particular pattern:
    - For example: starts with the letter a, ends with three uppercase letters, etc.
  - Why do these exist? Are the expressions the same across languages?
  - What's the difference between a basic and an extended regular expression? When would I use either?
  - How do I write a regular expression to accomplish x?
- 
- <http://regexpal.com/> to test your regex expressions
  - Simple regex tutorial [http://www.icewarp.com/support/online\\_help/203030104.htm](http://www.icewarp.com/support/online_help/203030104.htm)



# 4 Basic Concepts

- Quantification
  - How many times of previous expression?
  - Most common quantifiers: ?(0 or 1), \*(0 or more), +(1 or more)
- Grouping
  - Which subset of previous expression?
  - Grouping operator: ()
- Alternation
  - Which choices?
  - Operators: [] and |
    - Hello|World      [A B C]
- Anchors
  - Where?
  - Characters: ^ (beginning) and \$ (end)
- How do I use a combination of the above to accomplish tasks?

# Regular Expressions

Character	BRE / ERE	Meaning in a pattern
\	Both	Usually, turn off the special meaning of the following character. Occasionally, enable a special meaning for the following character, such as for <code>\(...\)</code> and <code>\{...\}</code> .
.	Both	Match any single character except NULL. Individual programs may also disallow matching newline.
*	Both	Match any number (or none) of the single character that immediately precedes it. For EREs, the preceding character can instead be a regular expression. For example, since <code>.</code> (dot) means any character, <code>.*</code> means "match any number of any character." For BREs, <code>*</code> is not special if it's the first character of a regular expression.
^	Both	Match the following regular expression at the beginning of the line or string. BRE: special only at the beginning of a regular expression. ERE: special everywhere.

# Regular Expressions (cont'd)

\$	Both	Match the preceding regular expression at the end of the line or string. BRE: special only at the end of a regular expression. ERE: special everywhere.
[...]	Both	Termed a bracket expression, this matches any one of the enclosed characters. A hyphen (-) indicates a range of consecutive characters. (Caution: ranges are locale-sensitive, and thus not portable.) A circumflex (^) as the first character in the brackets reverses the sense: it matches any one character not in the list. A hyphen or close bracket (]) as the first character is treated as a member of the list. All other metacharacters are treated as members of the list (i.e., literally). Bracket expressions may contain collating symbols, equivalence classes, and character classes (described shortly).
\{n,m\}	BRE	Termed an <i>interval expression</i> , this matches a range of occurrences of the single character that immediately precedes it. \{n\} matches exactly n occurrences, \{n,\} matches at least n occurrences, and \{n,m\} matches any number of occurrences between n and m. n and m must be between 0 and RE_DUP_MAX (minimum value: 255), inclusive.
\( \)	BRE	Save the pattern enclosed between \( and \) in a special <i>holding space</i> . Up to nine subpatterns can be saved on a single pattern. The text matched by the subpatterns can be reused later in the same pattern, by the escape sequences \1 to \9. For example, \(ab\).*\1 matches two occurrences of ab, with any number of characters in between.

# Regular Expressions (cont'd)

<code>\n</code>	BRE	Replay the nth subpattern enclosed in <code>\(</code> and <code>\)</code> into the pattern at this point. n is a number from 1 to 9, with 1 starting on the left.
<code>{n,m}</code>	ERE	Just like the BRE <code>\{n,m\}</code> earlier, but without the backslashes in front of the braces.
<b>+</b>	ERE	Match one or more instances of the preceding regular expression.
<b>?</b>	ERE	Match zero or one instances of the preceding regular expression.
<b> </b>	ERE	Match the regular expression specified before or after.
<b>( )</b>	ERE	Apply a match to the enclosed group of regular expressions.

# Examples

Expression	Matches
<code>tolstoy</code>	The seven letters tolstoy, anywhere on a line
<code>^tolstoy</code>	The seven letters tolstoy, at the beginning of a line
<code>tolstoy\$</code>	The seven letters tolstoy, at the end of a line
<code>^tolstoy\$</code>	A line containing exactly the seven letters tolstoy, and nothing else
<code>[Tt]olstoy</code>	Either the seven letters Tolstoy, or the seven letters tolstoy, anywhere on a line
<code>tol.toy</code>	The three letters tol, any character, and the three letters toy, anywhere on a line
<code>tol.*toy</code>	The three letters tol, any sequence of zero or more characters, and the three letters toy, anywhere on a line (e.g., tolstoy, tolWHOttoy, and so on)

# Text Processing Tools

- You should be familiar with:
  - `wc`: outputs a one-line report of lines, words, and bytes
  - `head`: extract top of files
  - `tail`: extracts bottom of files
  - `tr`: translate or delete characters
  - `grep`: print lines matching a pattern
  - `sort`: sort lines of text files
  - `sed`: filtering and transforming text
- What are the differences between `tr`, `sed`, and `grep`?  
When would I use each one?
- How can I combine and use these tools together?

# sort, comm, and tr

**sort**: sorts **lines** of **text** files

- Usage: sort [OPTION]...[FILE]...
- Sort order depends on locale
- C locale: ASCII sorting

**comm**: compare two **sorted** files **line by line**

- Usage: comm [OPTION]...FILE1 FILE2
- Comparison depends on locale

**tr**: translate **or** delete characters

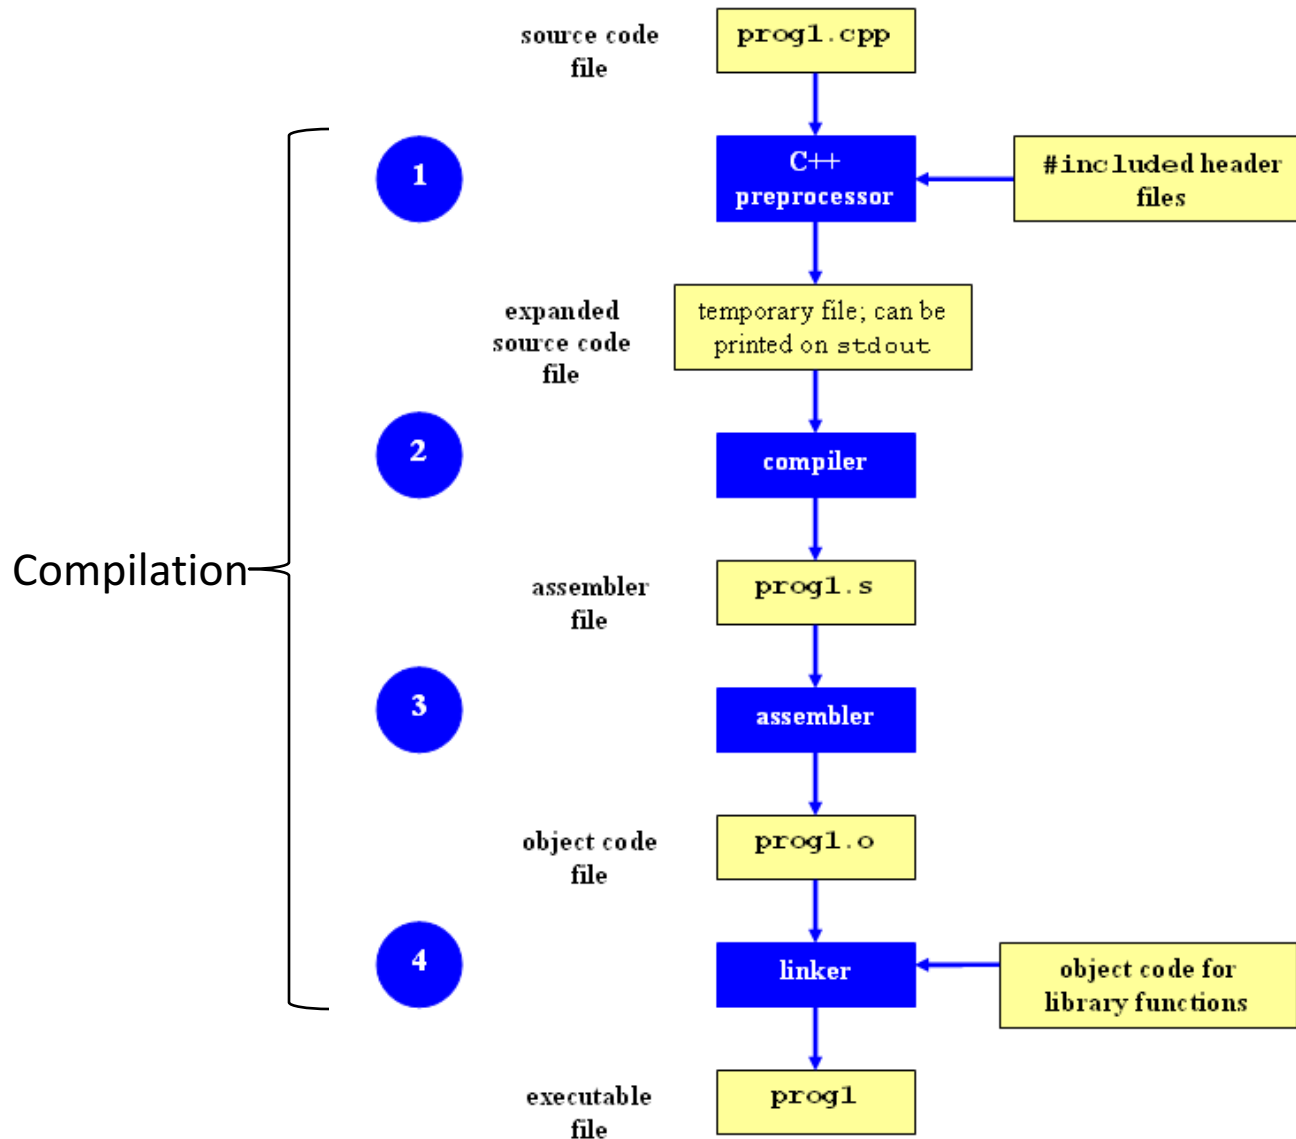
- Usage: tr [OPTION]...SET1 [SET2]

You've implemented comm and tr by hand, do you remember how you did that?

Week 3



# Compilation Process



# Compilation Process

- Why do we have this process?
- What are the different components of the process?
  - “I just typed gcc to compile my programs... does that mean gcc has all of the components within it?”
- Why can't I execute individual object code files?
- What are the differences between open source and closed source software? When would I want to use one or the other?

# Make

- Utility for managing large software projects
- Compiles files and keeps them up-to-date
- Efficient Compilation (only files that need to be recompiled)
- Why do we have make at all?

# Build Process

- **configure**
  - Script that checks details about the machine before installation
    - Dependency between packages
  - Creates 'Makefile'
- **make**
  - Requires 'Makefile' to run
  - Compiles all the program code and creates executables in current temporary directory
- **make install**
  - make utility searches for a label named install within the Makefile, and executes only that section of it
  - executables are copied into the final directories (system directories)

# Patching

- A patch is a piece of software designed to fix problems with or update a computer program
- It's a diff file that includes the changes made to a file
- A person who has the original (buggy) file can use the patch command with the diff file to add the changes to their original file
- Why not just change the original source code to fix it? Why do we have patches?

# Applying a Patch

Source Files



Original File

Modified File



Patch File



Original File



Patch File



Modified File

# diff Unified Format

- `diff -u original_file modified_file`
- `--- path/to/original_file`
- `+++ path/to/modified_file`
- `@@ -l,s +l,s @@`
  - `@@`: beginning of a hunk
  - `l`: beginning line number
  - `s`: number of lines the change hunk applies to for each file
  - A line with a:
    - `-` sign was deleted from the original
    - `+` sign was added to the original
    - stayed the same

# What is Python?

- Not just a scripting language
- Object-Oriented language
  - Classes
  - Member functions
- Compiled and interpreted
  - Python code is compiled to bytecode
  - Bytecode interpreted by Python interpreter
- Not as fast as C but easy to learn, read and use
- You should know how to write basic programs in python



# Comm.py

- Support all options for `comm`
  - -1, -2, -3 and combinations
  - Extra option `-u` for comparing unsorted files
- Support all type of arguments
  - File names and `-` for stdin
- Be familiar with how the linux `comm` utility works
- You should be able to write the `comm` utility by hand

Week 4

# Software development process

- Involves making a lot of changes to code
  - New features added
  - Bugs fixed
  - Performance enhancements
- Software team has many people working on the same/different parts of code
- Many versions of software released
  - Ubuntu 10, Ubuntu 12, etc
  - Need to be able to fix bugs for Ubuntu 10 for customers using it, even though you have shipped Ubuntu 12.

How do we deal with all of this?

# Source/Version Control

- Track changes to code and other files related to the software
  - What new files were added?
  - What changes made to files?
  - Which version had what changes?
  - Which user made the changes?
- Track entire history of the software
- Version control software
  - GIT, Subversion, Perforce

This seems complicated. Why bother with source control?

What are the strengths and weaknesses of source control?

When would I want to use it? How do I use it?

# Terms used

- **Repository**
  - Files and folder related to the software code
  - Full History of the software
- **Working copy**
  - Copy of software's files in the repository
- **Check-out**
  - To create a working copy of the repository
- **Check-in / Commit**
  - Write the changes made in the working copy to the repository
  - Commits are recorded by the VCS

# Terms used

- Head
  - Refers to a commit object
  - There can be many heads in a repository
- HEAD
  - Refers to the currently active head
- Detached HEAD
  - If a commit is not pointed to by a branch
  - This is okay if you want to just take a look at the code and if you don't commit any new changes
  - If the new commits have to be preserved then a new branch has to be created
    - `git checkout v3.0 -b BranchVersion3.1`
- Branch
  - Refers to a head and its entire set of ancestor commits
- Master
  - Default branch

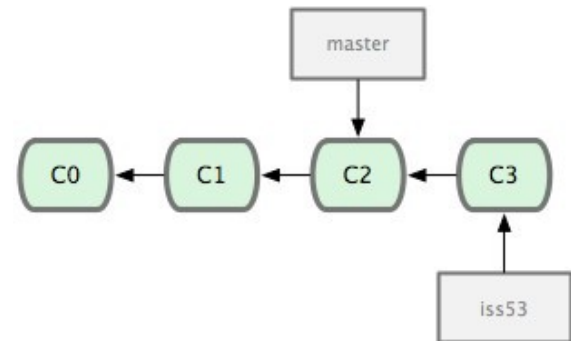


Image Source: git-scm.com

# What Is a Branch?

- A pointer to one of the commits in the repo (head) + all ancestor commits
- When you first create a repo, are there any branches?
  - Default branch named 'master'
- The default master branch
  - points to last commit made
  - moves forward automatically, every time you commit

# Questions

- What is the difference between a working copy and the repository?
- What is a commit? What should be in a commit? How many files should commits contain?
- What are the differences between head and HEAD?
- Why bother having branches at all? Why can't we just all work on the same single master branch?
- What happens when we perform a merge? How does it work?



# Git States

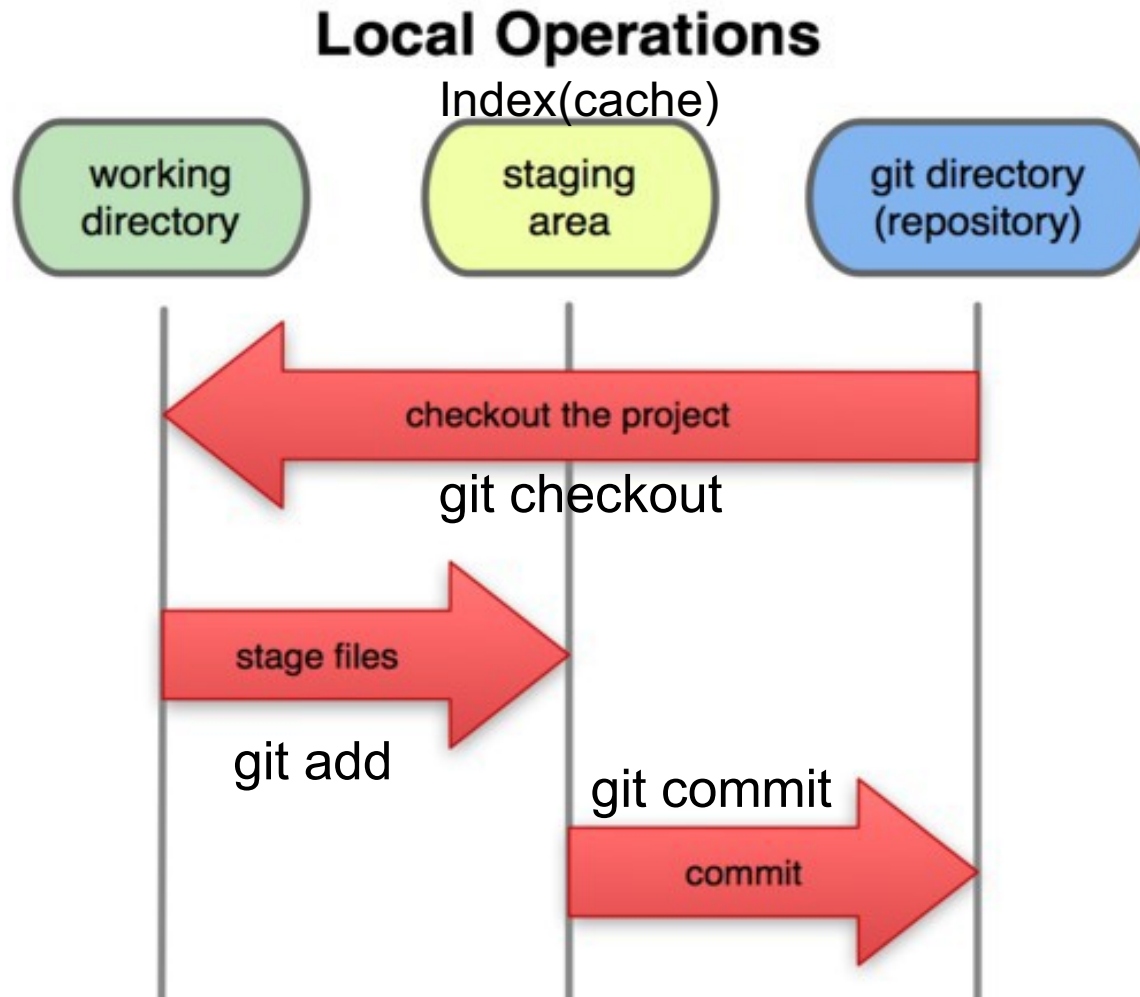


Image Source: [git-scm.com](http://git-scm.com)

# Git commands

- Repository creation
  - `$ git init` (Start a new repository)
  - `$ git clone` (Create a copy of an existing repository)
- Branching
  - `$ git checkout <tag/commit> -b <new_branch_name>` (creates a new branch)
- Commits
  - `$ git add` (Stage modified/new files)
  - `$ git commit` (check-in the changes to the repository)
- Getting info
  - `$ git status` (Shows modified files, new files, etc)
  - `$ git diff` (compares working copy with staged files)
  - `$ git log` (Shows history of commits)
  - `$ git show` (Show a certain object in the repository)
- Getting help
  - `$ git help`

You should be familiar with how these commands work and when to use them.

# More Git Commands

- Reverting
  - `$ git checkout HEAD main.cpp`
    - Gets the HEAD revision for the working copy
  - `$ git checkout -- main.cpp`
    - Reverts changes in the working directory
  - `$ git revert`
    - Reverting commits (this creates new commits)
- Cleaning up untracked files
  - `$ git clean`
- Tagging
  - Human readable pointers to specific commits
  - `$ git tag -a v1.0 -m 'Version 1.0'`
    - This will name the HEAD commit as v1.0

You should be familiar with how these commands work and when to use them.

Week 5

# Debugger

- A program that is used to run and debug other (target) programs
- Advantages:
  - Programmer can:
    - step through source code line by line
      - each line is executed on demand
    - interact with and inspect program at run-time
    - If program crashes, the debugger outputs where and why it crashed
- Why have a debugger?
- When do I use a debugger?

# Using GDB

## 1. Compile Program

- Normally: `$ gcc [flags] <source files> -o <output file>`
- Debugging: `$ gcc [other flags] -g <source files> -o <output file>`
  - enables built-in debugging support

## 2. Specify Program to Debug

- `$ gdb <executable>`
- `$ gdb`
- `(gdb) file <executable>`

# Using GDB

## 3. Run Program

- `(gdb) run`                      or
- `(gdb) run [arguments]`

## 4. In GDB Interactive Shell

- Tab to Autocomplete, up-down arrows to recall history
- `help [command]` to get more info about a command

## 5. Exit the gdb Debugger

- `(gdb) quit`

# Run-Time Errors

- Segmentation fault
  - Program received signal SIGSEGV, Segmentation fault.  
0x0000000000400524 in *function* (arr=0x7fffc902a270, r1=2, c1=5, r2=4, c2=6) at *file.c*:12
    - Line number where it crashed and parameters to the function that caused the error
- Logic Error
  - Program will run and exit successfully
- How do we find bugs?



# Setting Breakpoints

- Breakpoints
  - used to stop the running program at a specific point
  - If the program reaches that location when running, it will pause and prompt you for another command
- Example:
  - (gdb) break file1.c:6
    - Program will pause when it reaches line 6 of file1.c
  - (gdb) break my\_function
    - Program will pause at the first line of `my_function` every time it is called
  - (gdb) break [*position*] if *expression*
    - Program will pause at specified position only when the expression evaluates to true
- How do we know where to set breakpoints?
- What do we do once we've stopped at a breakpoint?

# Deleting, Disabling and Ignoring BPs

- (gdb) delete [bp\_number | range]
  - Deletes the specified breakpoint or range of breakpoints
- (gdb) disable [ *bp\_number* | *range* ]
  - Temporarily deactivates a breakpoint or a range of breakpoints
- (gdb) enable [ *bp\_number* | *range* ]
  - Restores disabled breakpoints
- If no arguments are provided to the above commands, all breakpoints are affected!!
- (gdb) ignore *bp\_number iterations*
  - Instructs GDB to pass over a breakpoint without stopping a certain number of times.
    - bp\_number: the number of a breakpoint
    - Iterations: the number of times you want it to be passed over

# Displaying Data

- Why would we want to interrupt execution?
  - to see data of interest at run-time:
  - (gdb) `print [/format] expression`
    - Prints the value of the specified expression in the specified format
  - Formats:
    - d: Decimal notation (default format for integers)
    - x: Hexadecimal notation
    - o: Octal notation
    - t: Binary notation
- What's the point of displaying data?
- What sort of data might we want to display?
- How can we use displayed data?

# Resuming Execution After a Break

- When a program stops at a breakpoint
  - 4 possible kinds of gdb operations:
    - **c or continue**: debugger will continue executing until next breakpoint
    - **s or step**: debugger will continue to next source line
    - **n or next**: debugger will continue to next source line in the current (innermost) stack frame
    - **f or finish**: debugger will resume execution until the current function returns. Execution stops immediately after the program flow returns to the function's caller
      - the function's return value and the line containing the next statement are displayed
- What is the difference between 's' and 'n'?
- When would we use each one of the above?

# Stack Info

- A program is made up of one or more functions which interact by calling each other
- Every time a function is called, an area of memory is set aside for it. This area of memory is called a **stack frame** and holds the following crucial info:
  - storage space for all the local variables
  - the memory address to return to when the called function returns
  - the arguments, or parameters, of the called function
- Each function call gets its own stack frame. Collectively, all the stack frames make up the **call stack**
- Why does the stack exist at all? How is the stack different than the heap?

# Analyzing the Stack in GDB

- `(gdb) backtrace | bt`
  - Shows the call trace (the call stack)
  - Without function calls:
    - #0 main () at program.c:10
    - one frame on the stack, numbered 0, and it belongs to main()
  - After call to function `display()`
    - #0 display (z=5, zptr=0xbffffb34) at program.c:15
    - #1 0x08048455 in main () at program.c:10
    - Two stack frames: frame 1 belonging to main() and frame 0 belonging to display().
    - Each frame listing gives
      - the arguments to that function
      - the line number that's currently being executed within that frame

# C Programming

- You should be able to develop a basic C program that incorporates the following:
  - Basic data types
  - Control/flow (if, while, etc)
  - Pointers
  - Structs
  - Dynamic memory
  - Basic I/O

# Dynamic Memory

- Memory that is allocated at runtime
  - Why?
- Allocated on the heap
  - Why not the stack?

**void \*malloc (size\_t size);**

- Allocates *size* bytes and returns a pointer to the allocated memory

**void \*realloc (void \*ptr, size\_t size);**

- Changes the size of the memory block pointed to by *ptr* to *size* bytes

**void free (void \*ptr);**

- Frees the block of memory pointed to by *ptr*

- What happens if I never call free?
- What happens if I try to put data into dynamic memory but I haven't yet called malloc?

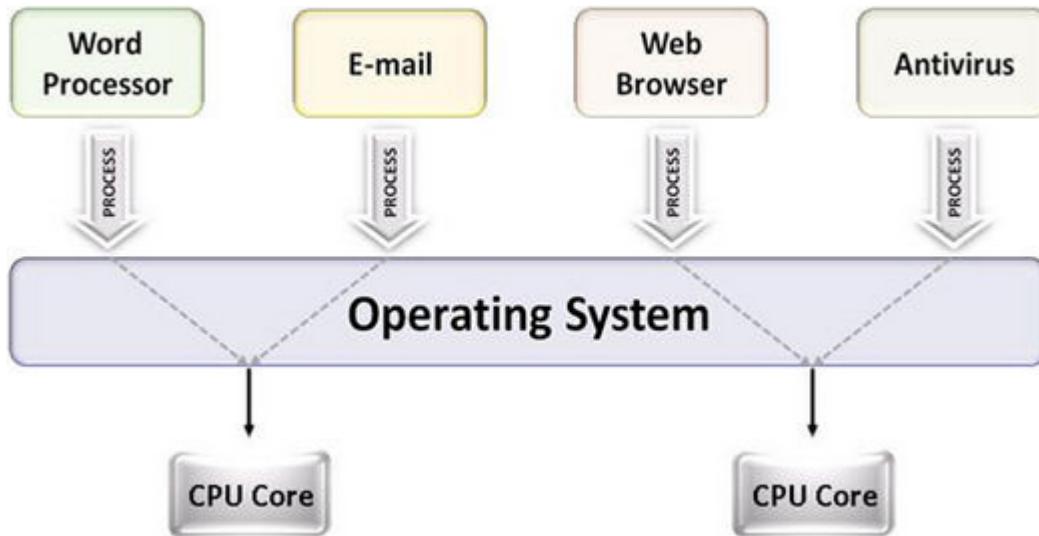


Week 6

# Parallelism

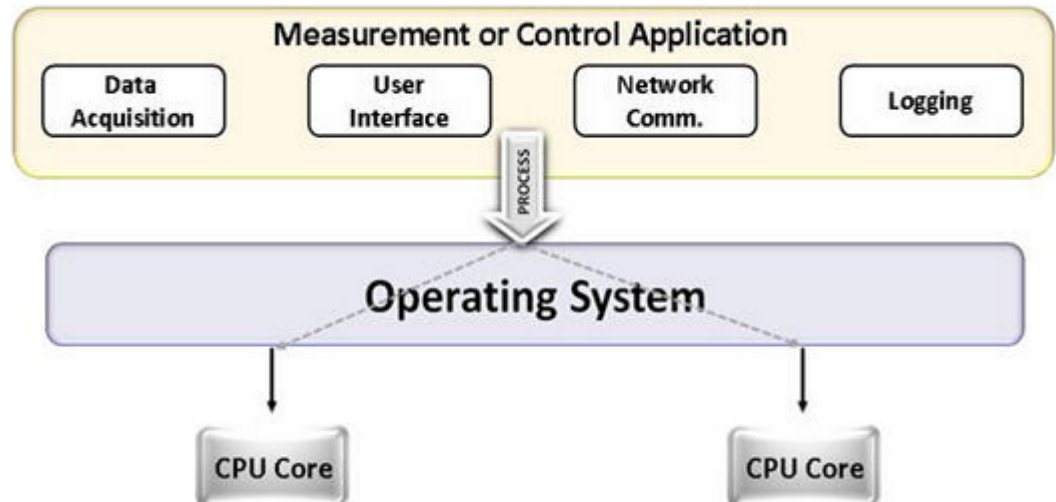
- Executing several computations simultaneously to gain performance
- Different forms of parallelism
  - **Multitasking**
    - Several processes are scheduled alternately or possibly simultaneously on a multiprocessing system
  - **Multithreading**
    - Same job is broken logically into pieces (threads) which may be executed simultaneously on a multiprocessing system
- What's the point of parallelism? Isn't it just too complicated?
- How can you decide whether your application should use multiple processes or multiple threads? Or both?

# Multitasking vs. Multithreading



**Multitasking**

**Multithreading**



# Multithreading & Multitasking: Comparison

- **Multithreading**

- Threads share the same address space
  - Light-weight creation/destruction
  - Easy inter-thread communication
  - An error in one thread can bring down all threads in process

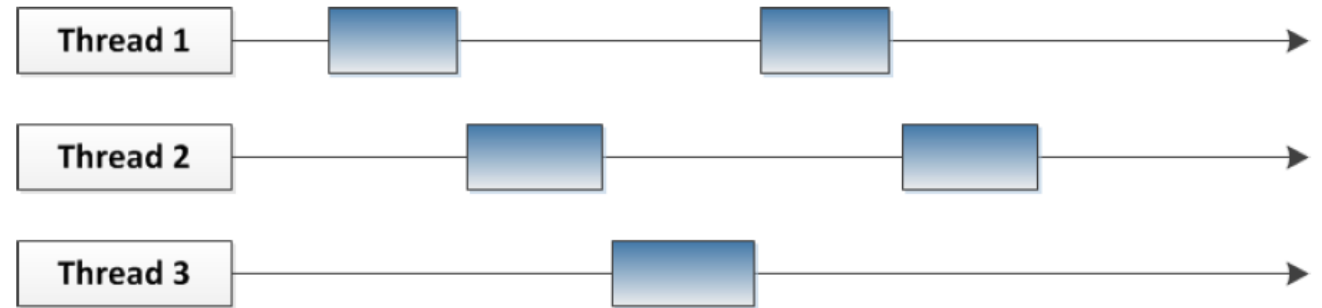
- **Multitasking**

- Processes are insulated from each other
  - Expensive creation/destruction
  - Expensive IPC
  - An error in one process cannot bring down another process

# What is a thread?

- A flow of instructions, path of execution within a process
- The smallest unit of processing scheduled by OS
- A process consists of at least one thread
- Multiple threads can be run on:
  - **A uniprocessor (time-sharing)**
    - Processor switches between different threads
    - Parallelism is an illusion
  - **A multiprocessor**
    - Multiple processors or cores run the threads at the same time
    - True parallelism

Multiple threads sharing a single CPU



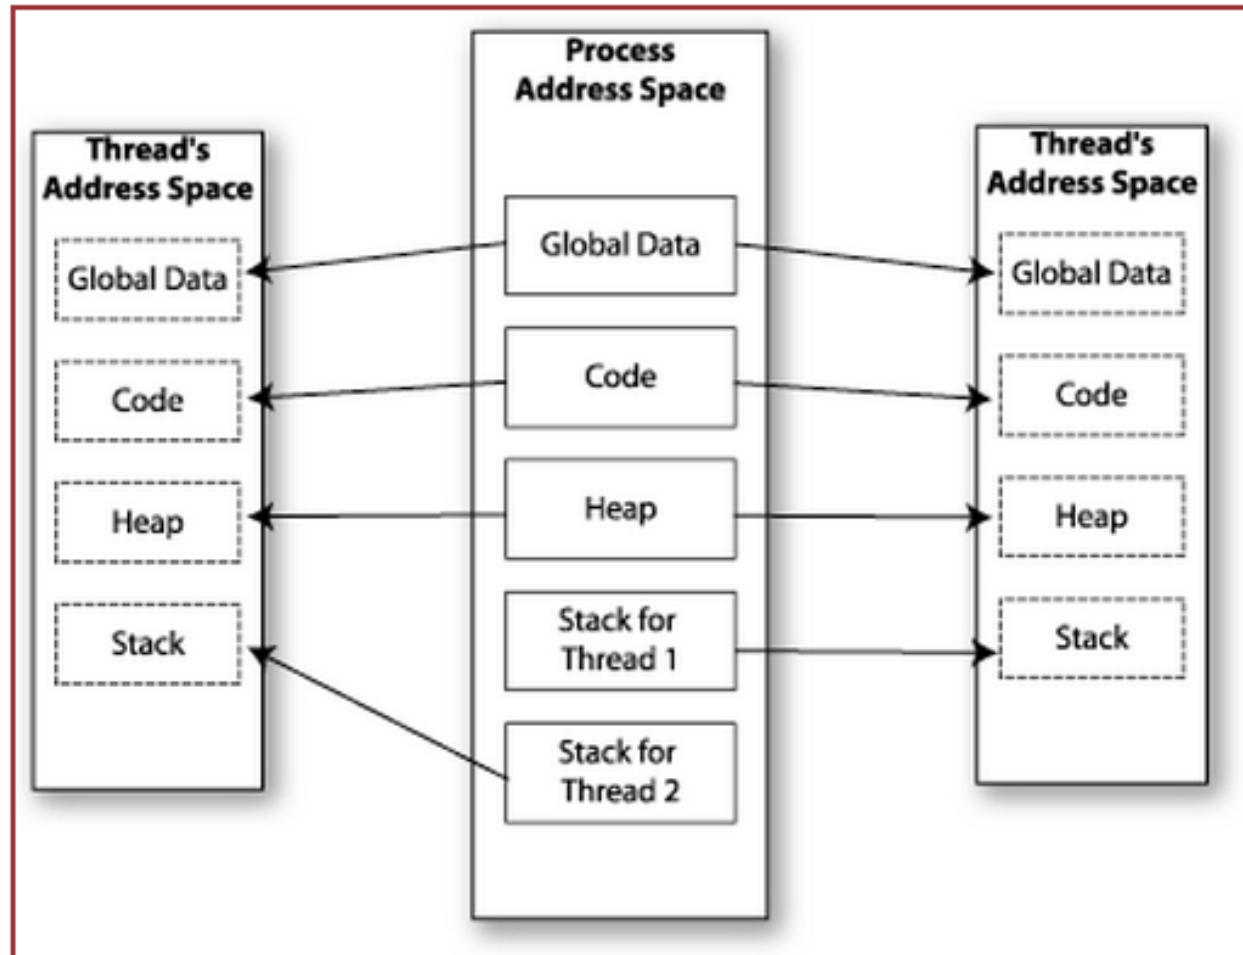
Multiple threads on multiple CPUs



# Memory Layout: Multithreaded Program



Why share the  
same code?  
Why share the  
same heap?



# Shared Memory

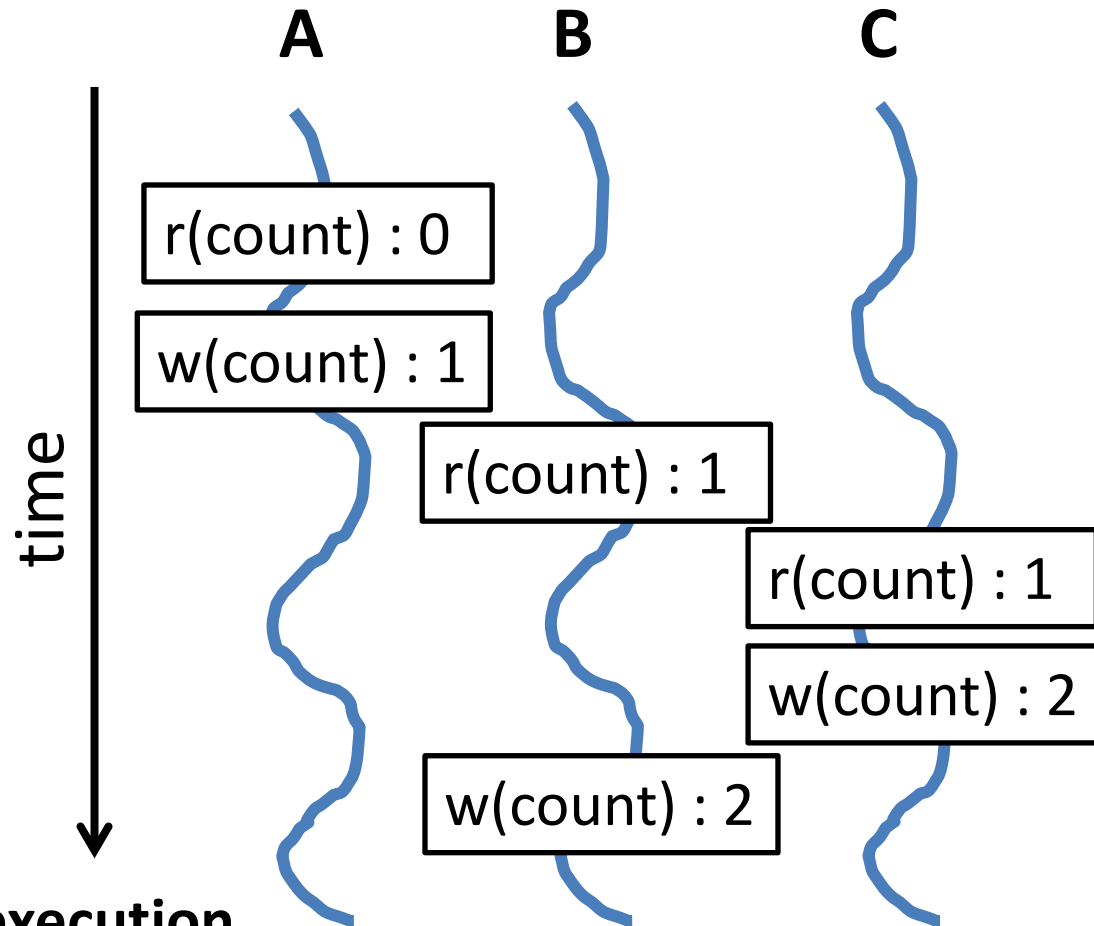
- Makes multithreaded programming
  - **Powerful**
    - can easily access data and share it among threads
  - **More efficient**
    - No need for system calls when sharing data
    - Thread creation and destruction less expensive than process creation and destruction
  - **Non-trivial**
    - Have to prevent several threads from accessing and changing the same shared data at the same time (synchronization)

What happens when we have race conditions?



# Race Condition

```
int count = 0;  
void increment()  
{  
    count = count + 1;  
}
```



Result depends on order of execution  
=> Synchronization needed

# pthread\_create

- **Function:** creates a new thread and makes it executable
- Can be called any number of times from anywhere within code
- Return value:
  - Success: zero
  - Failure: error number
- How do we keep track of threads within a program's execution? How many can we have?
- How do we pass data to threads we create? How do we tell them what to work on?
- What happens if our application isn't "embarrassingly parallel"?

# Parameters

```
int pthread_create( pthread_t *tid, const pthread_attr_t *attr,  
                  void *(my_function)(void *), void *arg );
```

- **tid**: unique identifier for newly created thread
- **attr**: object that holds thread attributes (priority, stack size, etc.)
  - Pass in NULL for default attributes
- **my\_function**: function that thread will execute once it is created
- **arg**: a *single* argument that may be passed to my\_function
  - Pass in NULL if no arguments

# pthread\_join

- **Function:** makes originating thread wait for the completion of all its spawned threads' tasks
- Without join, the originating thread would exit as soon as it completes its job
  - ⇒ A spawned thread can get aborted even if it is in the middle of its chore
- Return value:
  - Success: zero
  - Failure: error number
- Why join at all? What does a join guarantee?

# Arguments

```
int pthread_join(pthread_t tid, void **status);
```

- **tid**: thread ID of thread to wait on
- **status**: the exit status of the target thread is stored in the location pointed to by \*status
  - Pass in NULL if no status is needed

Week 7

# Processor Modes

- Operating modes that place restrictions on the type of operations that can be performed by running processes
  - User mode: restricted access to system resources
  - Kernel/Supervisor mode: unrestricted access
- System resources?
  - Memory
  - I/O Devices
  - CPU
- Why have different modes? How do we switch modes?

# User Mode vs. Kernel Mode

- Hardware contains a mode-bit, e.g. 0 means kernel mode, 1 means user mode
- User mode
  - CPU **restricted** to unprivileged instructions and a specified area of memory
- Supervisor/kernel mode
  - CPU is **unrestricted**, can use all instructions, access all areas of memory and take over the CPU anytime
- What happens if user code is given unrestricted access to CPU?



# Why Dual-Mode Operation?

- System resources are shared among processes
- OS must ensure:
  - **Protection**
    - an incorrect/malicious program cannot cause damage to other processes or the system as a whole
  - **Fairness**
    - Make sure processes have a fair use of devices and the CPU

# How to Achieve Protection and Fairness

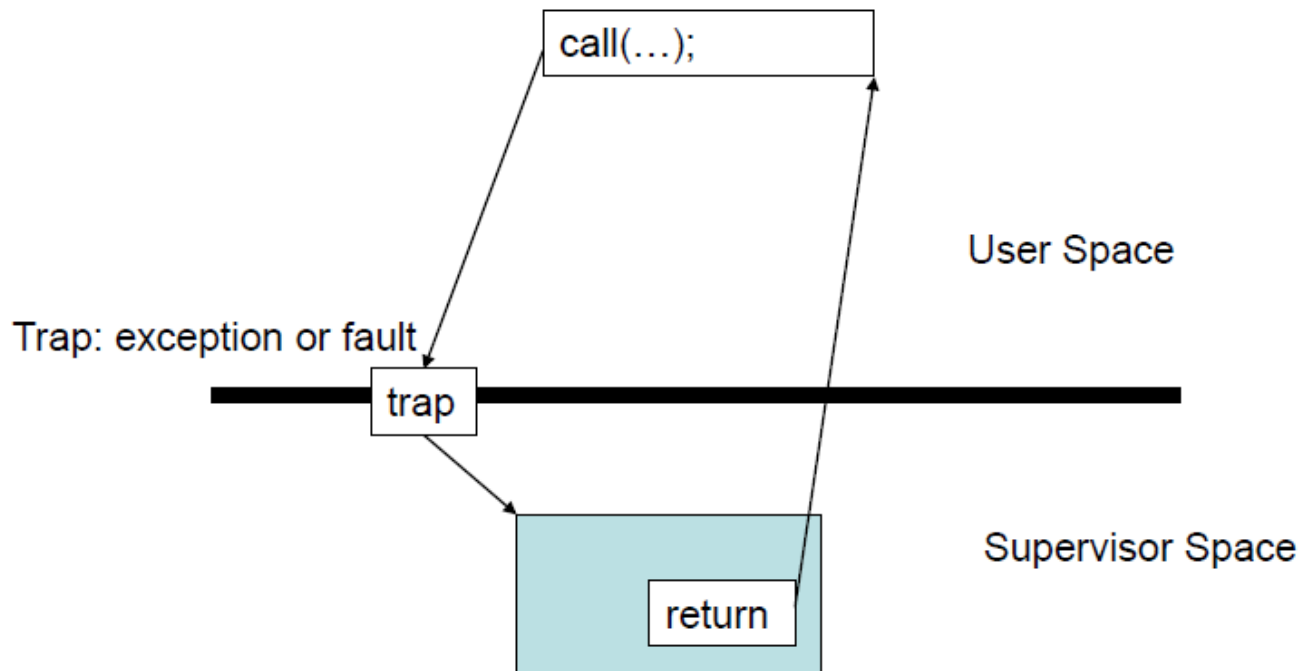
- Goals:
    - **I/O Protection**
      - Prevent processes from performing illegal I/O operations
    - **Memory Protection**
      - Prevent processes from accessing illegal memory and modifying kernel code and data structures
    - **CPU Protection**
      - Prevent a process from using the CPU for too long
- => instructions that might affect goals are privileged and can only be executed by *trusted code*

# System Calls

- Special type of function that:
  - Used by user-level processes to request a service from the kernel
  - Changes the CPU's mode from user mode to kernel mode to enable more capabilities
  - Is part of the kernel of the OS
  - Verifies that the user should be allowed to do the requested action and then does the action (kernel performs the operation on behalf of the user)
  - Is the ***only way*** a user program can perform privileged operations
- When do I need to use system calls?

# System Calls

- When a system call is made, the program being executed is interrupted and control is passed to the kernel
- If operation is valid the kernel performs it



# System Call Overhead

- System calls are expensive and can hurt performance
- The system must do many things
  - Process is interrupted & computer saves its state
  - OS takes control of CPU & verifies validity of op.
  - **OS performs requested action**
  - OS restores saved context, switches to user mode
  - OS gives control of the CPU back to user process

# Library Functions

- Functions that are a part of standard C library
- To avoid system call overhead use equivalent library functions
  - getchar, putchar vs. read, write (for standard I/O)
  - fopen, fclose vs. open, close (for file I/O), etc.
- How do these functions perform privileged operations?
  - They make system calls
- What are the benefits and tradeoffs of using either system calls or C library functions?

# Unbuffered vs. Buffered I/O

- **Unbuffered**

- Every byte is read/written by the kernel through a system call

- **Buffered**

- collect as many bytes as possible (in a buffer) and read more than a single byte (into buffer) at a time and use one system call for a block of bytes

⇒ Buffered I/O decreases the number of read/write system calls and the corresponding overhead

Which is faster in what applications? When would you use buffered or unbuffered I/O?

Week 8



# Static Linking

- Carried out only once to produce an executable file
- If static libraries are called, the linker will copy all the modules referenced by the program to the executable
- Static libraries are typically denoted by the .a file extension
- When would I use static linking? Why would I use it?

# Dynamic Linking

- Allows a process to add, remove, replace or relocate object modules during its execution.
- If shared libraries are called:
  - Only copy a little reference information when the executable file is created
  - Complete the linking during loading time or running time
- Dynamic libraries are typically denoted by the .so file extension
  - .dll on Windows
- When would I use dynamic linking? Why would I use it?

# Linking and Loading

- Linker collects procedures and links them together object modules into one executable program
- Why isn't everything written as just one **big** program, saving the necessity of linking?
  - Efficiency: if just one function is changed in a 100K line program, why recompile the whole program? Just recompile the one function and relink.
  - Multiple-language programs
  - Other reasons?
- When does linking happen? When does loading happen?

# How are libraries dynamically loaded?

**Table 1. The DL API**

Function	Description
<b>dlopen</b>	Makes an object file accessible to a program
<b>dlsym</b>	Obtains the address of a symbol within a dlopened object file
<b>dlerror</b>	Returns a string error of the last error that occurred
<b>dlclose</b>	Closes an object file

# Dynamic linking

- Unix systems: Code is typically compiled as a *dynamic shared object* (DSO)
- Dynamic vs. static linking resulting size

```
$ gcc -static hello.c -o hello-static
$ gcc hello.c -o hello-dynamic
$ ls -l hello
    80 hello.c
 13724 hello-dynamic
   383 hello.s
1688756 hello-static
```
- If you are the sysadmin, which do you prefer?
- What is the difference between linking and loading?

# Advantages of dynamic linking

- The executable is typically smaller
- When the library is changed, the code that references it does not usually need to be recompiled
- The executable accesses the .so at run time; therefore, multiple programs can access the same .so at the same time
  - Memory footprint amortized across all programs using the same .so
- What other advantages are there of dynamic linking?

# Disadvantages of dynamic linking

- Performance hit
  - Need to load shared objects (at least once)
  - Need to resolve addresses (once or every time)
  - Remember back to the system call assignment...
- What if the necessary dynamic library is missing?
- What if we have the library, but it is the wrong version?

Week 9



# Communication Over the Internet

- What type of guarantees do we want?
  - **Confidentiality**
    - Message secrecy
  - **Data integrity**
    - Message consistency
  - **Authentication**
    - Identity confirmation
  - **Authorization**
    - Specifying access rights to resources
- Why do we want these guarantees?

# Cryptography

- **Plaintext** – Actual message
- **Ciphertext** – Encrypted message (unreadable gibberish)
- **Encryption** – Going from plaintext to ciphertext
- **Decryption** – Going from ciphertext to plaintext
- **Secret key**
  - Part of the mathematical function used to encrypt/decrypt.
  - Good key makes it hard to get back plaintext from ciphertext

Be familiar with all of these terms and what they mean/represent.



Image Source: [gpgtools.org](http://gpgtools.org)

# Symmetric-key Encryption

- Same secret key used for encryption and decryption
- **Example** : Data Encryption Standard (**DES**)
- **Caesar's cipher**
  - Map the alphabet to a shifted version
    - ABCDEFGHIJKLMNOPQRSTUVWXYZ
    - DEF...GHIJKLMNOPQRSTUVWXYZABC
  - Plaintext – SECRET. Ciphertext – VHFUHW
  - Key is 3 (number of shifts of the alphabet)
- **Key distribution** is a problem
  - The secret key has to be delivered in a safe way to the recipient
  - Chance of key being compromised

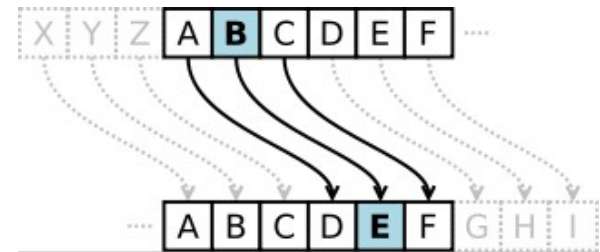


Image Source: wikipedia

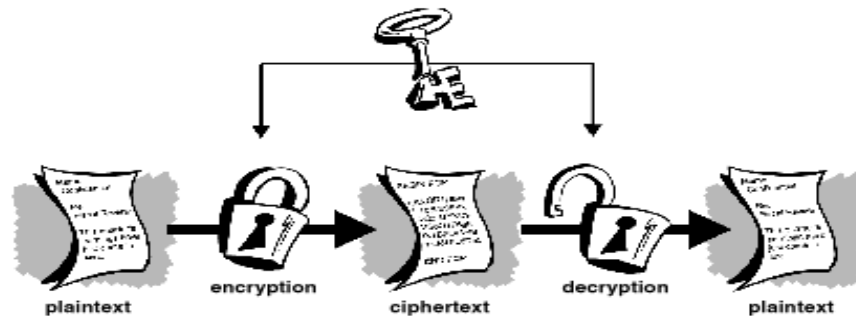
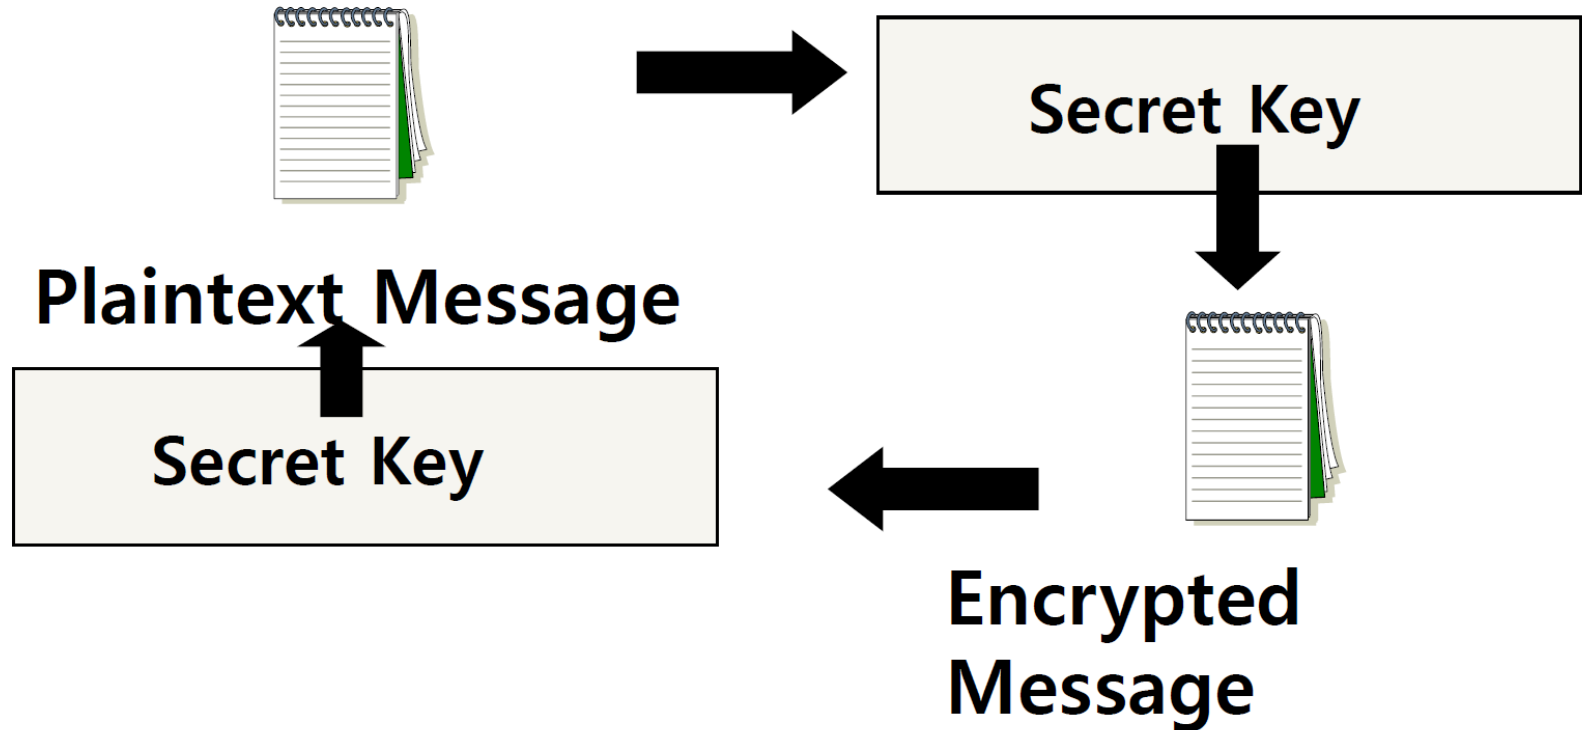


Image Source: gpgtools.org

# Secret Key (symmetric) Cryptography

- A single key is used to both encrypt and decrypt a message



# Public-key Encryption (Asymmetric)

- Uses a pair of keys for encryption
  - **Public key** – Published and known to everyone
  - **Private key** – Secret key known only to the owner
- **Encryption**
  - Use public key to encrypt messages
  - Anyone can encrypt message, but they cannot decrypt the ciphertext
- **Decryption**
  - Use private key to decrypt messages
- **Example : RSA** – Rivest, Shamir & Adleman
  - Property used - **Difficulty of factoring** large integers to prime numbers
  - $N = p * q$  (3233 = 61 \* 53)
  - N is a large integer and p, q are prime numbers
  - N is part of the public key
  - [http://en.wikipedia.org/wiki/RSA\\_Factoring\\_Challenge](http://en.wikipedia.org/wiki/RSA_Factoring_Challenge)

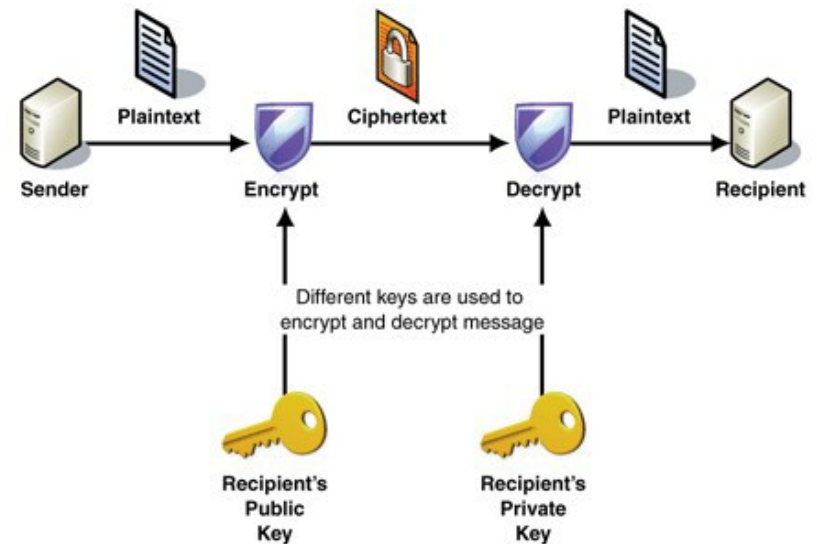
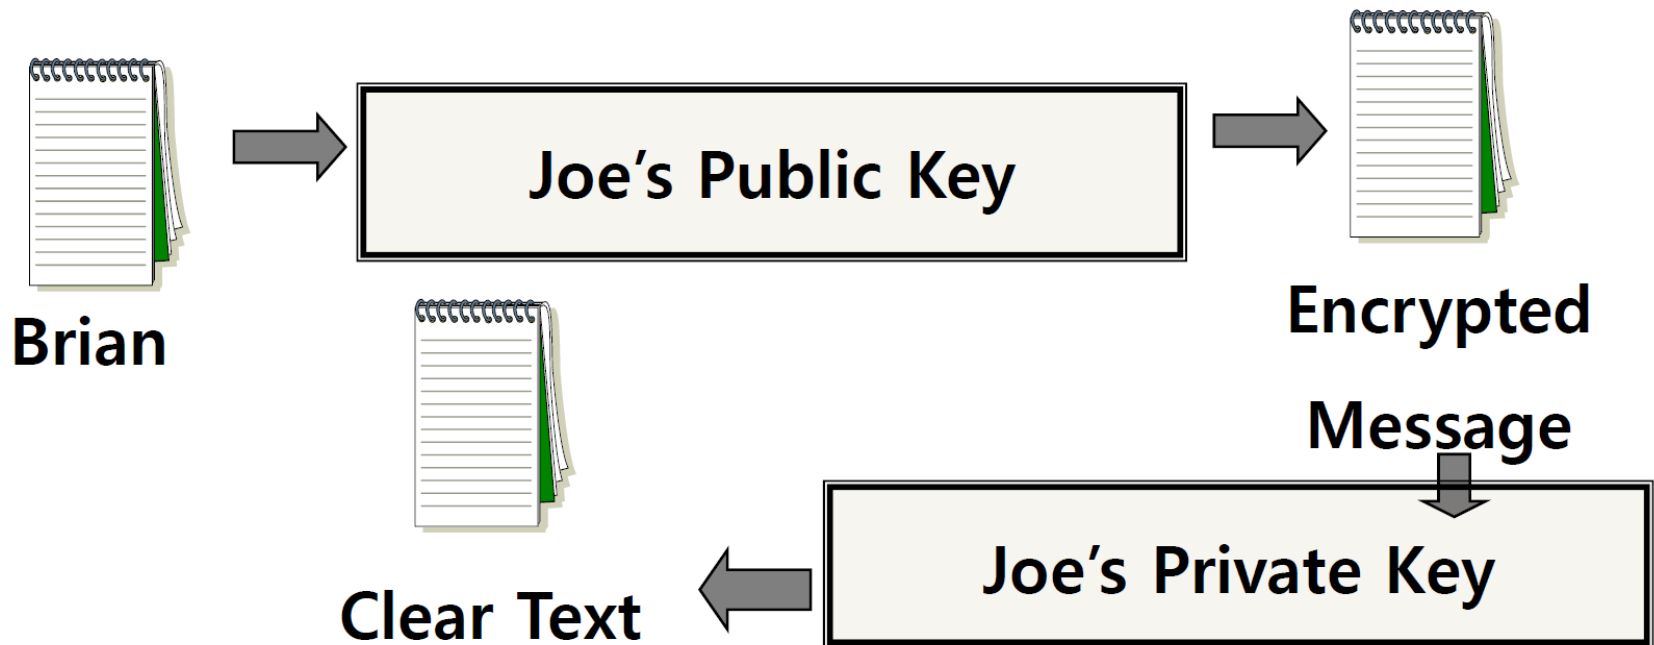


Image Source: MSDN

# Public Key (asymmetric) Cryptography

- Two keys are used: a public and a private key. If a message is encrypted with one key, it has to be decrypted with the other.



# Encryption questions

- Why have encryption?
- What are the differences between symmetric and asymmetric encryption? When would I use one or the other?
- What is used on the Internet? What is a certificate authority?
- How can I trust a message came from someone?

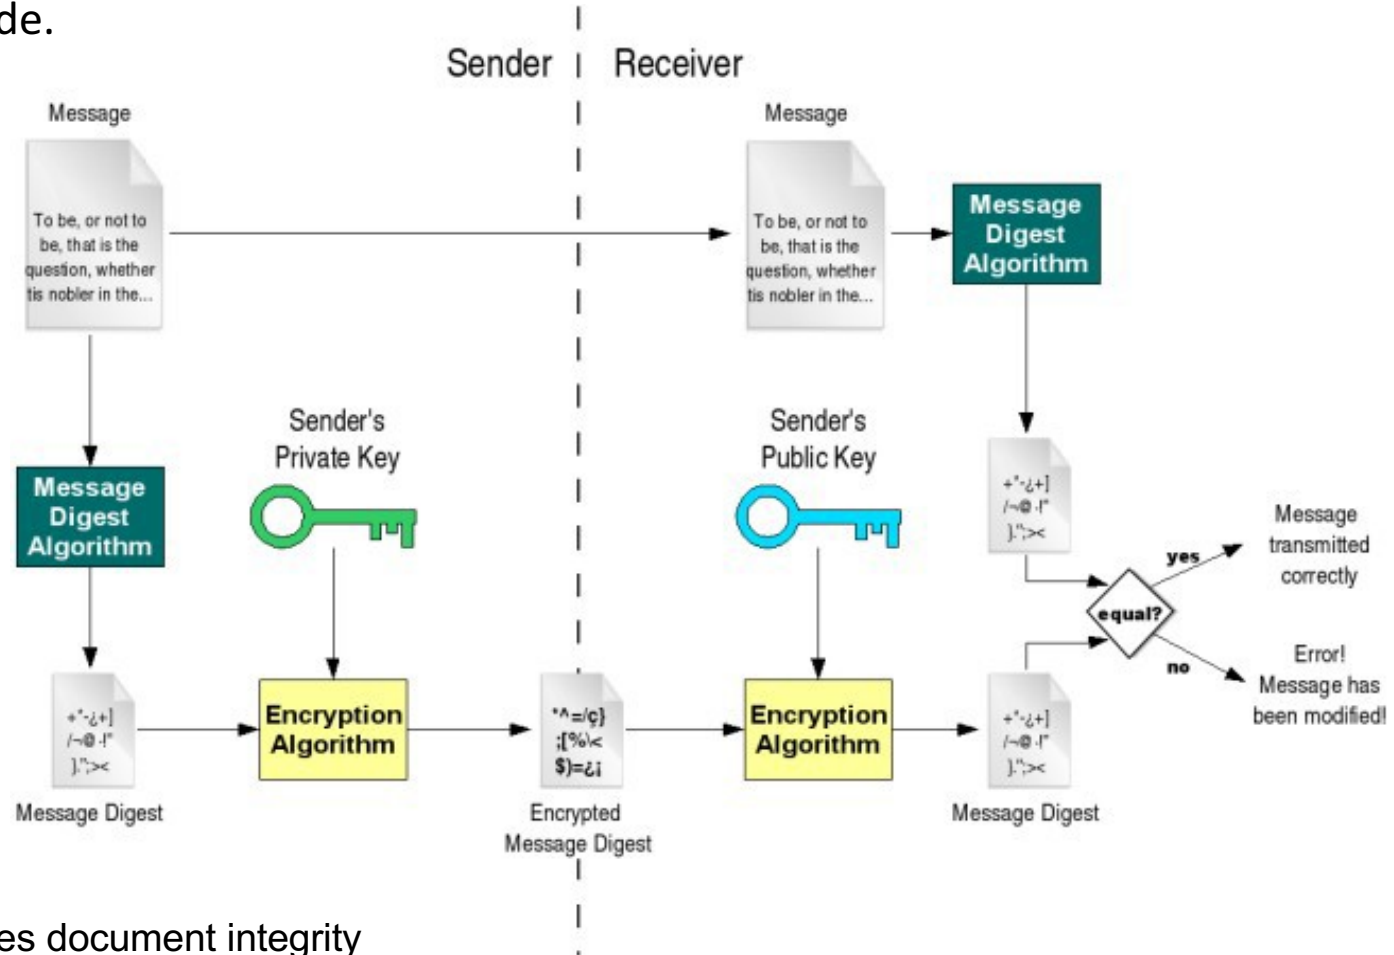
# Digital Signature

- An electronic stamp or seal
  - almost exactly like a written signature, except more guarantees!
- Is appended to a document
  - Or sent separately (detached signature)
- Ensures data integrity
  - document was not changed during transmission
- How are signatures different than encryption?



# Digital Signature

You should understand this entire slide.



- Verifies document integrity
- Does it prove origin?
- Who is Certificate Authority (CA) ?

Image Source : gdp.globus.org

# Detached Signature

- Digital signatures can either be *attached* to the message or *detached*
- A detached signature is stored and transmitted separately from the message it signs
- Commonly used to validate software distributed in compressed tar files
- You can't sign such a file internally without altering its contents, so the signature is created in a separate file
- Why detach at all? Why are signatures useful?
- Who can create a signature? How do I verify a signature?

Good luck for your finals!